
Go in TiDB

Yao Wei | PingCAP

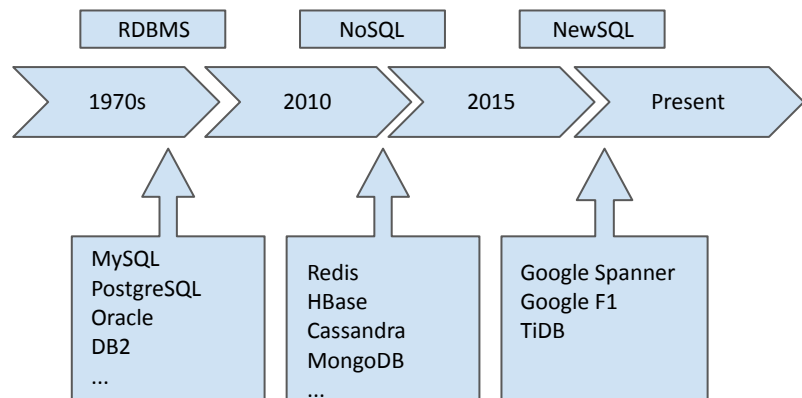
About me

- Yao Wei (姚维)
- TiDB Kernel Expert, General Manager of South Region, China
- 360 Infra team / Alibaba-UC / PingCAP
- Atlas/MySQL-Sniffer
- Infrastructure software engineer

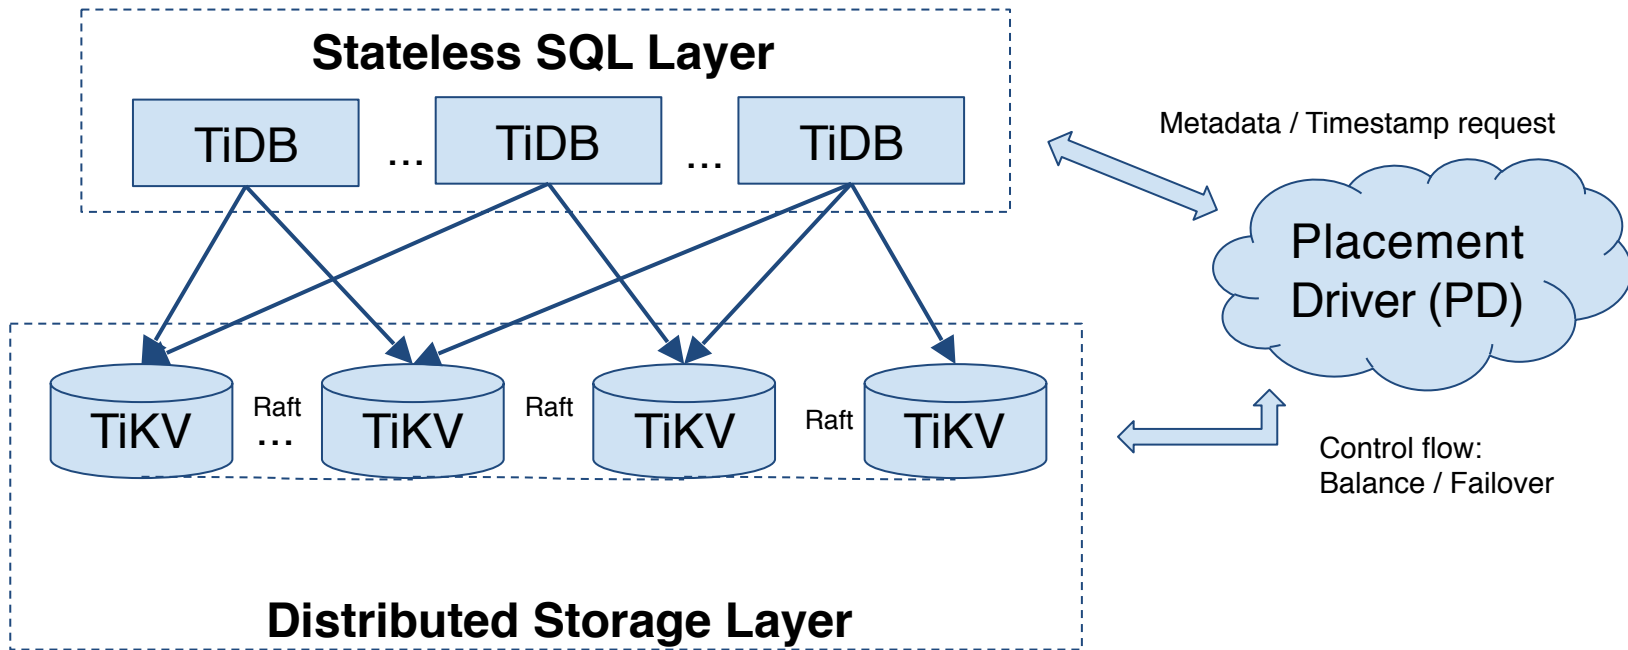
Why a new database?

Brief History

- Standalone RDBMS
- NoSQL
- Middleware & Proxy
- NewSQL





Architecture



TiKV - Overview

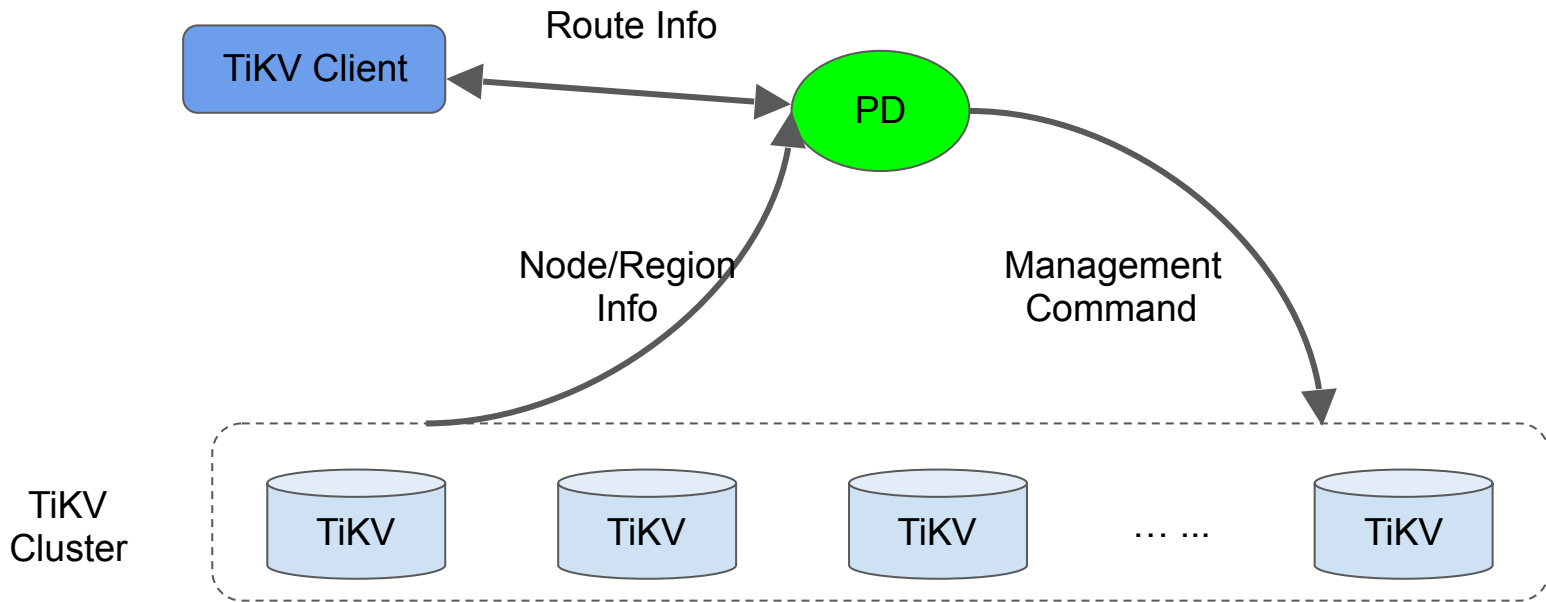
- Region: a set of continuous key-value pairs
- Data is organized/stored/replicated by Regions

- Highly leveraged



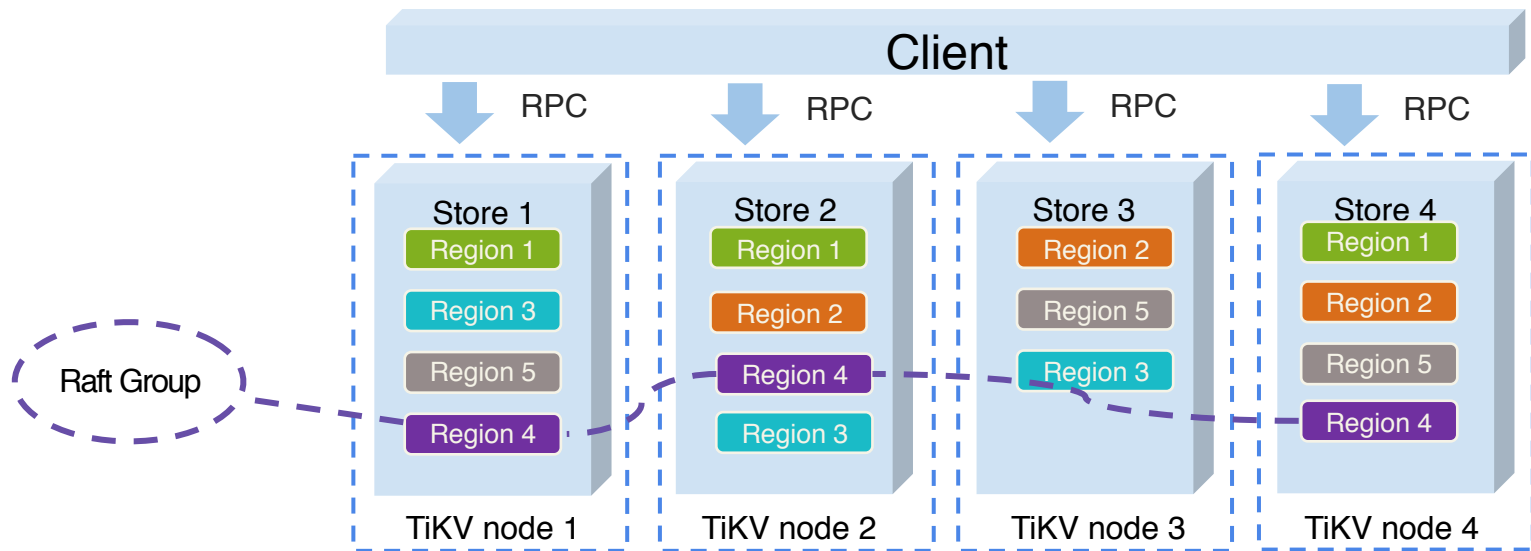
PD - Overview

- Meta data management
- Load balance management




TiKV - Multi-Raft

Multiple raft groups in the cluster, one group for each region.



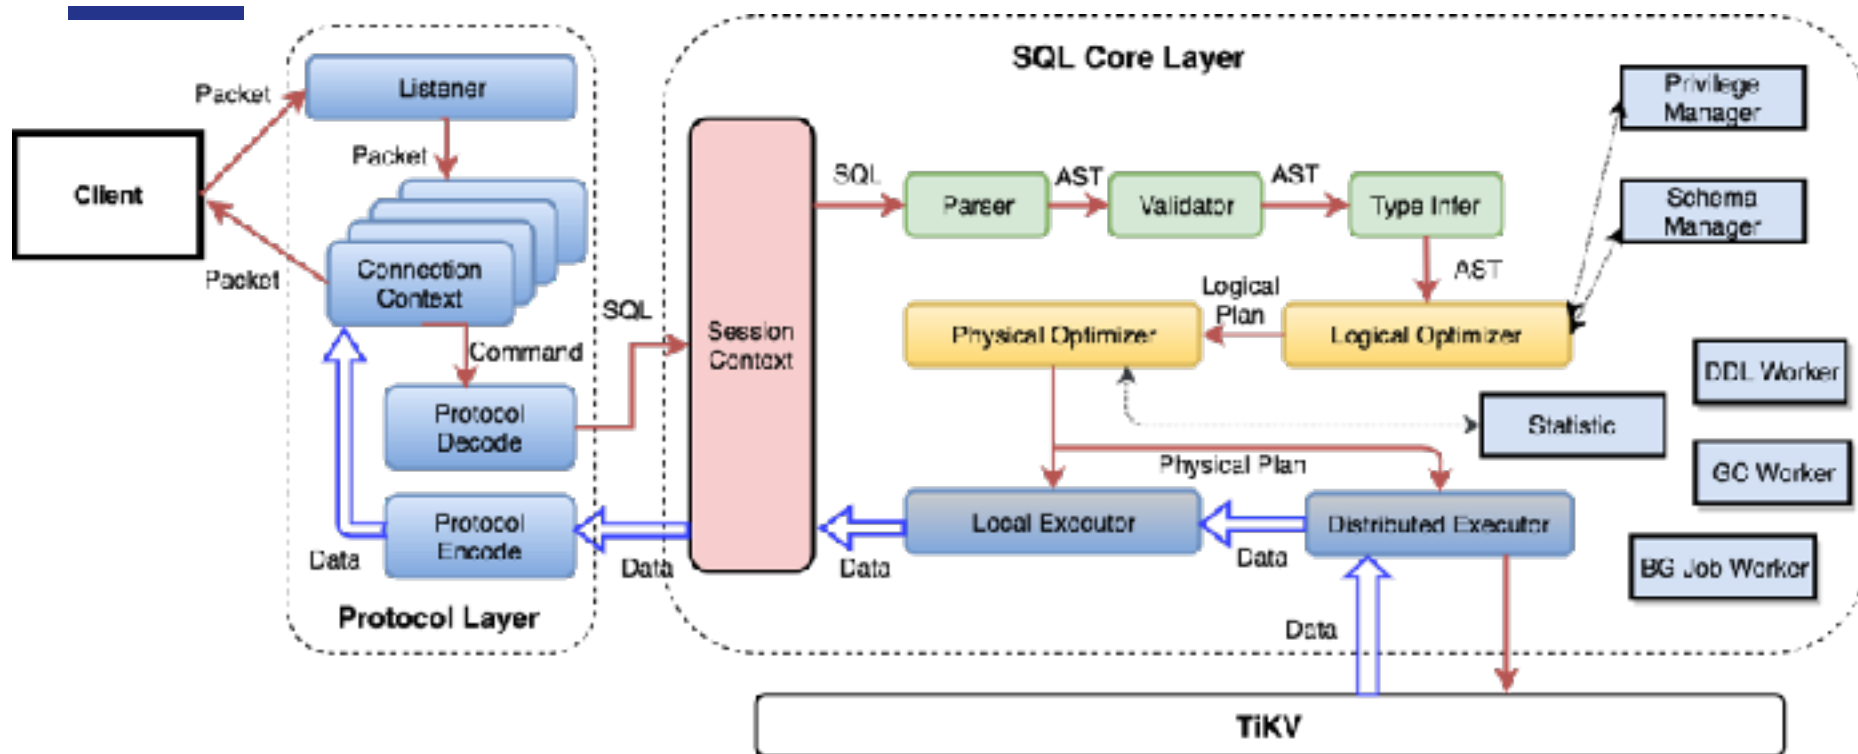
TiKV - Horizontal Scale



Three steps to move a leader replica

- Transfer Leader
- Add Replica
- Remove Replica

SQL Layer

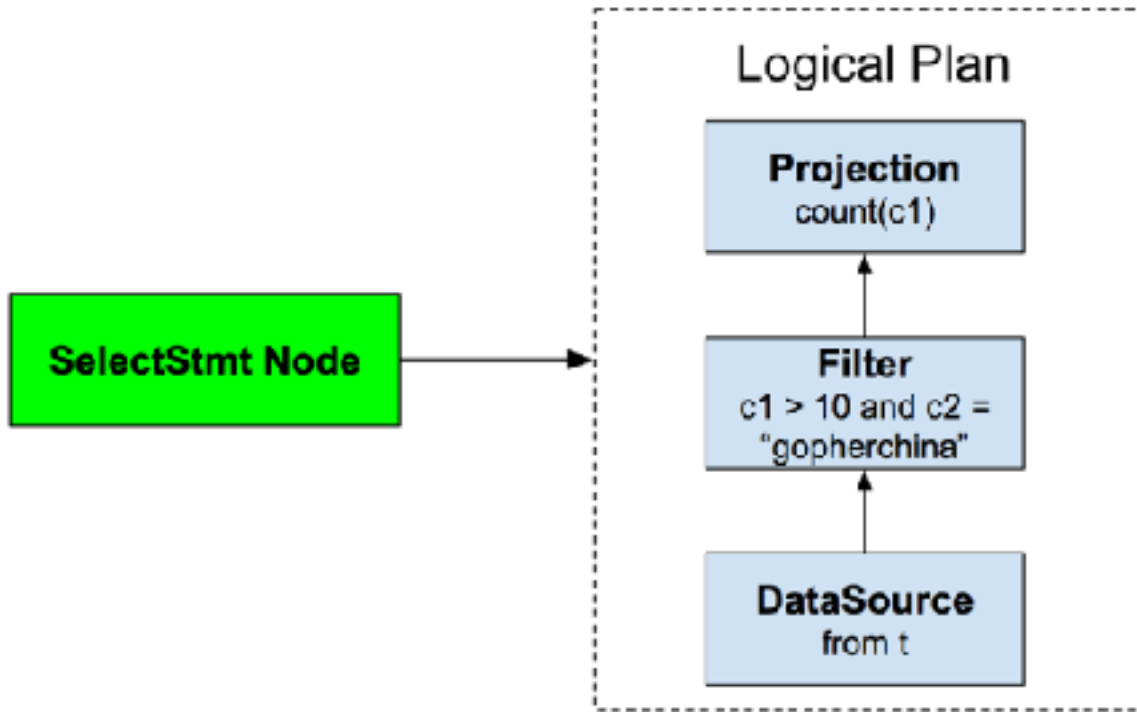


Example - SQL

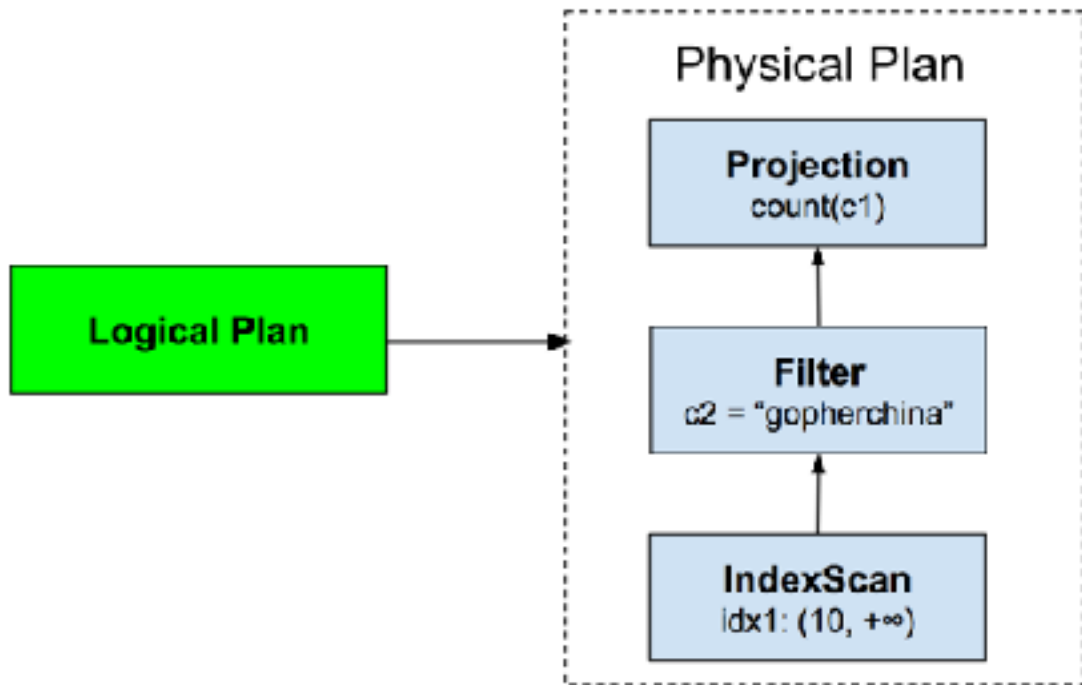
```
CREATE TABLE t (c1 INT, c2 VARCHAR(32), INDEX idx1 (c1));
```

```
SELECT COUNT(c1) FROM t WHERE c1 > 10 AND c2 =  
"gopherchina";
```

Example - Logical Plan



Example - Physical Plan



Challenges of distributed ACID database?

- Distributed Database is very complex
- Lots of RPC work
- Keep high performance
- Tons of data
- Huge amount of OLTP queries
- Very complex OLAP queries
- External Consistency
- SQL is much more complex than KV



Why TiDB choose Golang?

- Easy-learning
- Productivity
- Concurrency
- Easy to trace bugs and profile
- Standard libraries and tools
- Tolerant GC latency
- Good performance
- Quick improvement

Go in TiDB

- More than 160K lines of Go code and 138 contributors.

github.com/AIDanial/cloc v 1.72 T=4.00 s (137.4 files/s, 52274.8 lines/s)

Language	files	blank	comment	code
Go	523	17655	18761	163173
yacc	1	389	220	6236
XML	7	0	0	999
JSON	1	0	0	502
Markdown	9	238	0	470
YAML	3	6	4	227
make	1	38	2	132
Bourne Shell	3	12	5	52
Bourne Again Shell	1	9	18	47
Dockerfile	1	5	0	8
SUM:	550	18352	19010	171846

Memory && GC

- Query may touch a huge number of data.
- Memory allocation may cost a lot of time.
- Put pressure on GC worker.
- Degrade the performance of SQL.
- OOM sucks!
- `runtime.morestack`

Reduce the Number of Allocation

- Get enough memory in one allocation operation

```
a := []int{1, 2, 3, 4, 5}
```

```
b := []int{}
```

```
// a much better way:
```

```
// b := make([]int, 0, len(a))
```

```
for _, i := range a {
```

```
    b = append(b, i)
```

```
}
```

Reuse Object

- Share a stack for all queries in one session
- Introduce a cache in goyacc
- Resource pool

Reduce runtime.morestack

```
var growStack = false
```

```
func growStack() {  
    var buf [16 << 10] /* 16 KB */ byte  
    if growStack {  
        for i := range buf {  
            buf[i] = byte(i)  
        }  
        groupStack = true  
    }  
}
```

Not enough!

Goroutine Pool

[tidb/util/goroutine_pool/gp.go](#):

```
15 package gp
16
17 import (
18     "sync"
19     "sync/atomic"
20     "time"
21 )
22
23 // Pool is a struct to represent goroutine pool.
24 type Pool struct {
25     head    goroutine
26     tail    *goroutine
27     count    int
28     idleTimeout time.Duration
29     sync.Mutex
30 }
31
32 // goroutine is actually a background goroutine, with a channel binded for communication.
33 type goroutine struct {
34     ch    chan func()
35     next *goroutine
36     status int32
37 }
38
```

Row Memory Format

Id (Int)	name(varchar)	score(double)	salary(bigInt)
1	"a"	1.0	10000
2	"b"	1.2	20000
3	"c"	5.1	30000
4	"d"	7.9	4000

Row Memory Format

- How to store the row in memory?
- Union? Json? Protobuf?
- Use Datum to store one Column.

```
type Datum struct {  
    k      byte      // datum kind.  
    collation uint8    // collation can hold uint8 values.  
    decimal uint16    // decimal can hold uint16 values.  
    length  uint32    // length can hold uint32 values.  
    i      int64      // i can hold int64 uint64 float64 values.  
    b      []byte     // b can hold string or []byte values.  
    x      interface{} // x hold all other types.  
}
```

Row Memory Format

- **What is the disadvantage about Datum?**

- Use unnecessary memory in every column.
- Must use assert to get complex types:

```
func (d *Datum) GetMysqlDecimal() *MyDecimal {  
    return d.x.(*MyDecimal)  
}
```

- Impossible to do vectorizable serial computation
- So how to store multi types data in golang?

Apache Arrow

- binary data format
- Array lengths
- Null count
- Null bitmaps
- Offsets buffer
- Values Array
- more details plz see the docs

```

+ Length: 4, Null count: 1
+ Null bitmap buffer:

| Byte 0 (validity bitmap) | Bytes 1-63 |
|-----|-----|
| 00001001 | 0 (padding) |

+ Children arrays:
+ field-0 array (List<char>):
+ Length: 4, Null count: 2
+ Null bitmap buffer:

| Byte 0 (validity bitmap) | Bytes 1-63 |
|-----|-----|
| 00001000 | 0 (padding) |

+ Offsets buffer:

| Bytes 8-15 |
|-----|
| 0, 3, 3, 3, 7 |

+ Values array:
+ Length: 7, Null count: 0
+ Null bitmap buffer: Not required

+ Value buffer:

| Bytes 0-6 |
|-----|
| jansack |

+ field-1 array (Int32 array):
+ Length: 4, Null count: 1
+ Null bitmap buffer:

| Byte 0 (validity bitmap) | Bytes 1-63 |
|-----|-----|
| 00001001 | 0 (padding) |

+ Value buffer:

| Bytes 8-9 | | Bytes 10-11 | | Bytes 12-15 | | Bytes 16-23 | | |
|-----| |-----| |-----| |-----|
| 1 | | 2 | | unspecified | | 4 | | unspecified |

```

Row Memory Format

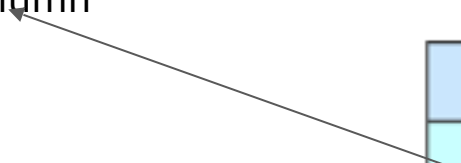
- We can eliminate the arrow's offsets array.
- Column store VS Row store.

id (int)	name(varchar)	score(double)	salary(bigint)
1	"a"	1.0	10000
2	"b"	1.2	20000
3	"c"	5.1	30000
4	"d"	7.9	4000

Chunk

```
type Chunk struct {  
    columns []*column  
}
```

```
type column struct {  
    length    int  
    nullCount int  
    nullBitmap []byte  
    offsets   []int32  
    data      []byte  
    elemBuf   []byte  
    ifaces    []interface{  
}
```



Id (Int)	name(varchar)	score(double)	salary(bigInt)
1	"a"	1.0	10000
2	"b"	1.2	20000
3	"c"	5.1	30000
4	"d"	7.9	4000

Chunk

- Use `unsafe.pointers` to get complex types:

```
func (r Row) GetMyDecimal(colIdx int) *types.MyDecimal {  
    col := r.c.columns[colIdx]  
    return (*types.MyDecimal)(unsafe.Pointer(&col.data[r.idx*types.MyDecimalStructSize]))  
}
```

```
func (r Row) GetUint64(colIdx int) uint64 {  
    col := r.c.columns[colIdx]  
    return *(*uint64)(unsafe.Pointer(&col.data[r.idx*8]))  
}
```

Chunk

- **Vectorized Execute expressions:**

```
func VectorizedExecute(ctx context.Context, exprs []Expression, input, output *chunk.Chunk) error {
    sc := ctx.GetSessionVars().StmtCtx
    for colID, expr := range exprs {
        err := evalOneColumn(sc, expr, input, output, colID)
        if err != nil {
            return errors.Trace(err)
        }
    }
    return nil
}

func evalOneColumn(sc *stmtctx.StatementContext, expr Expression, input, output *chunk.Chunk, colID int) (err error) {
    switch fieldType, evalType := expr.GetType(), expr.GetType().EvalType(); evalType {
    case types.ETInt:
        for row := input.Begin(); err == nil && row != input.End(); row = row.Next() {
            err = executeToInt(sc, expr, fieldType, row, output, colID)
        }
        ....
    }
}
```


Thanks!

Contact me:

wink@pingcap.com

www.pingcap.com

Wechat: 82091309



12.16 深圳 Gopher meetup



该二维码7天内(12月22日前)有效, 重新进入将更新